

XGBoost Regressor




What is XGBoost Regression?

- **XGBoost (Extreme Gradient Boosting)** is a powerful **machine learning algorithm** based on decision trees.
- It works by combining many small “weak” trees to build a **strong prediction model**.
- It is widely used because it is **fast, accurate, and works well on large datasets**.

Why XGBoost is Popular?

- ✓ Very fast (optimized for speed)
- ✓ Handles missing values automatically
- ✓ Prevents overfitting (regularization built-in)
- ✓ Works for **Regression & Classification**

XGBoost Regression

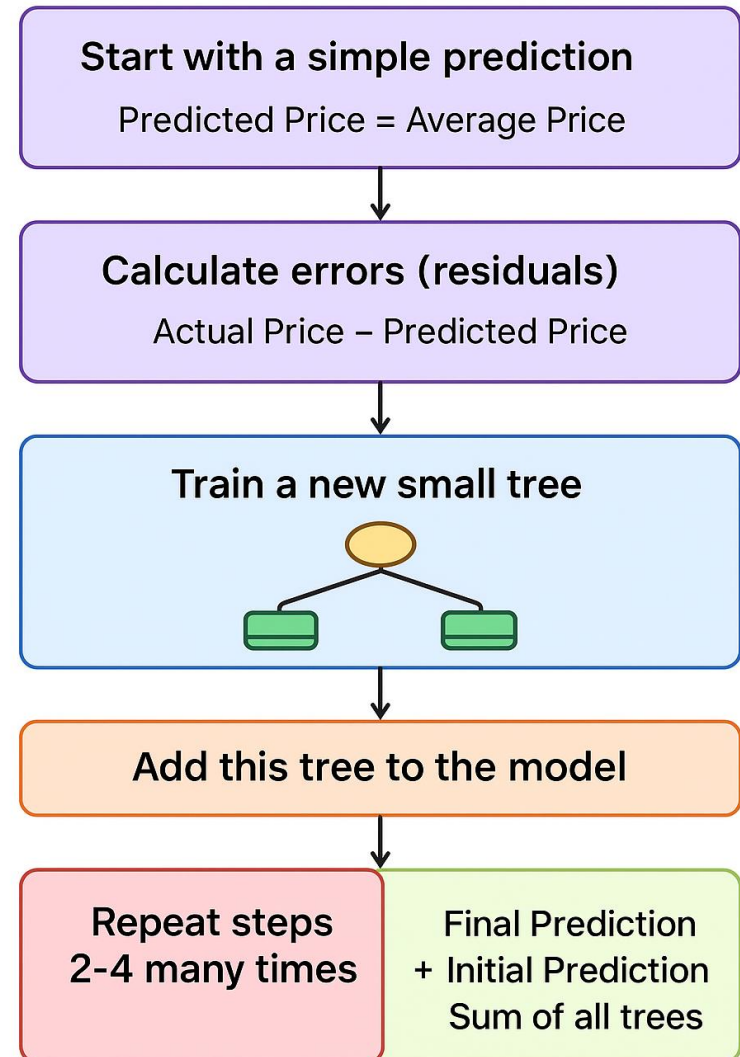


How It Works (Step-by-Step)



XGBoost Regressor

1. **Start with a simple prediction** (like predicting the average value for all houses).
2. **Calculate errors (residuals)** = Actual value – Predicted value.
3. **Train a new small tree** to fix those errors.
4. **Add this tree to the model** (with some weight).
5. Repeat steps 2–4 many times, each new tree fixing the mistakes of the previous ones.
6. Finally, combine all trees → gives a strong regressor.



Example: Predicting House Prices

Example: Predicting House Prices

Suppose we want to predict house price from **size (sqft)**, **location**, and **number of rooms**.

- Step 1: Model starts with an average prediction (e.g., ₹50 Lakhs).
- Step 2: Some houses are actually higher/lower → errors are calculated.
- Step 3: A new tree learns these errors (like: if size > 2000 sqft, add ₹20 Lakhs).
- Step 4: More trees are added, each fixing remaining mistakes.
- Step 5: After many rounds, the model gives very accurate house price predictions

Python

```
import xgboost as xg
```

```
model = xgb.XGBRegressor(  
    n_estimators=100,    # number of boosting rounds  
    learning_rate=0.1,  # step size shrinkage  
    max_depth=3,        # depth of trees  
    random_state=42  
)
```

Parameters:

.

n_estimators int, default=50

The maximum number of estimators at which boosting is terminated. In case of perfect fit, the learning procedure is stopped early. Values must be in the range `[1, inf)`.

random_state int, RandomState instance or None, default=None

Controls the random seed given at each `estimator` at each boosting iteration. Thus, it is only used when `estimator` exposes a `random_state`. In addition, it controls the bootstrap of the weights used to train the `estimator` at each boosting iteration. Pass an int for reproducible output across multiple function calls. See [Glossary](#).